

## Libraries and Imports

```
import pandas as pd
from pandas import DataFrame
from pandas import Series
import numpy as np
import colorsys
import cv2
import matplotlib.pyplot as plt
from numpy import ndarray
import itertools
import heapq
import time
import seaborn as sns
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
recall_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn import metrics
from sklearn.metrics import ConfusionMatrixDisplay
```

## Reading the data

```
aisles = pd.read_csv('aisles.csv')
departments = pd.read_csv('departments.csv')
products = pd.read_csv('products.csv')
orders = pd.read_csv('orders.csv')
target = pd.read_csv('target.csv')
order_products__prior = pd.read_csv('order_products__prior.csv')
order_products__train_test =
pd.read_csv('order_products__train_test.csv')
```

## Section A: Exploring the data

orders

	order_id	user_id	eval_set	order_number	order_dow	
0	2539329	1	prior	1	2	\
1	2398795	1	prior	2	3	
2	473747	1	prior	3	3	
3	2254736	1	prior	4	4	
4	431534	1	prior	5	4	
...	...	...	...	...	...	
3346078	2134068	124107	test	7	6	

3346079	2113602	70278	test	100	4
3346080	2342325	175026	test	4	4
3346081	3367040	179234	test	19	5
3346082	1828606	98632	test	4	4

	order_hour_of_day	days_since_prior_order
0	8	NaN
1	7	15.0
2	12	21.0
3	7	29.0
4	15	28.0
...	...	...
3346078	9	7.0
3346079	13	5.0
3346080	17	11.0
3346081	23	30.0
3346082	14	27.0

[3346083 rows x 7 columns]

departments.head()

	department_id	department
0	1	frozen
1	2	other
2	3	bakery
3	4	produce
4	5	alcohol

order\_products\_\_prior

	order_id	product_id	add_to_cart_order	reordered
0	2	33120	1	1
1	2	28985	2	1
2	2	9327	3	0
3	2	45918	4	1
4	2	30035	5	0
...	...	...	...	...
32434484	3421083	39678	6	1
32434485	3421083	11352	7	0
32434486	3421083	4600	8	0
32434487	3421083	24852	9	1
32434488	3421083	5020	10	1

[32434489 rows x 4 columns]

order\_products\_\_train\_test

	order_id	product_id	add_to_cart_order	reordered
0	1	49302	1	1
1	1	11109	2	1

2	1	10246	3	0
3	1	49683	4	0
4	1	43633	5	1
...	...	...	...	...
1384612	3421063	14233	3	1
1384613	3421063	35548	4	1
1384614	3421070	35951	1	1
1384615	3421070	16953	2	1
1384616	3421070	4724	3	1

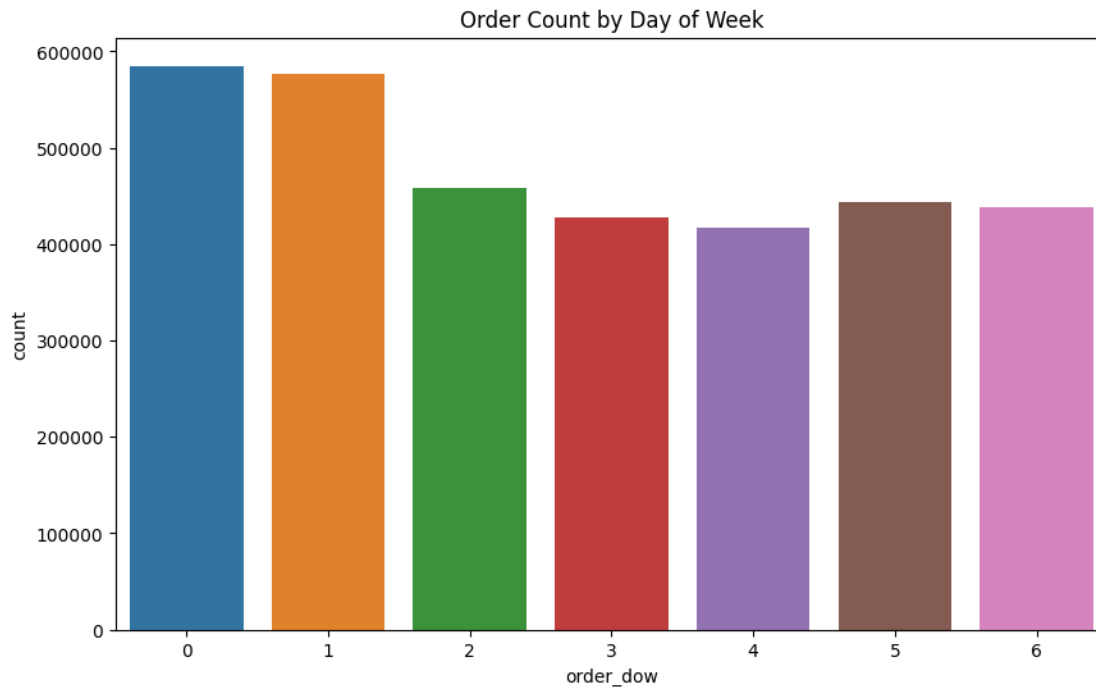
[1384617 rows x 4 columns]

target

	order_id	product_id	was_in_order
0	1187899	196.0	1
1	2757217	196.0	1
2	632715	196.0	1
3	1167274	196.0	1
4	3347074	196.0	1
...	...	...	...
2959120	3421070	4987.0	0
2959121	3421070	8230.0	0
2959122	3421070	39468.0	0
2959123	3421070	39139.0	0
2959124	3421070	48208.0	0


[2959125 rows x 3 columns]

```
plt.figure(figsize=(10,6))
sns.countplot(x='order_dow', data=orders)
plt.title('Order Count by Day of Week')
plt.show()
```



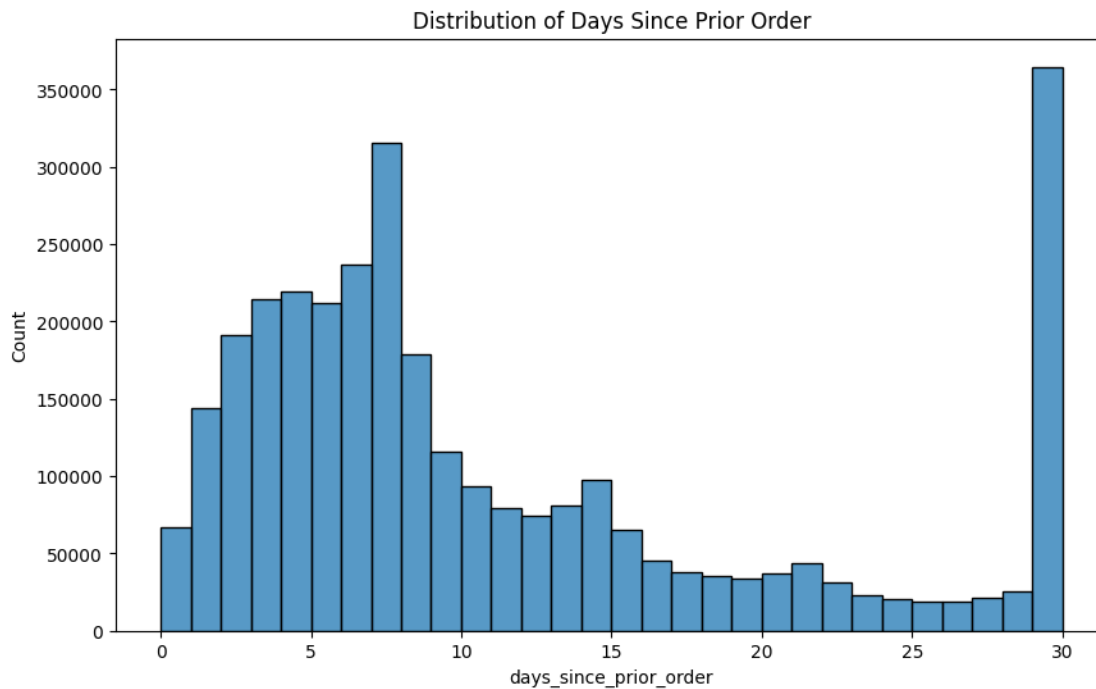
we can see that most people tend to make orders at the beginning of the week

```
plt.figure(figsize=(10,6))
sns.countplot(x='order_hour_of_day', data=orders)
plt.title('Order Count by Hour of the Day')
plt.show()
```



```
plt.figure(figsize=(10,6))
sns.histplot(orders['days_since_prior_order'].dropna(), bins=30,
```

```
kde=False)
plt.title('Distribution of Days Since Prior Order')
plt.show()
```




**We can say that users tend to place and order every 7 days, 30 days.**

```
reorders = order_products__prior['reordered'].value_counts()
```


```
plt.figure(figsize=(6,6))
plt.pie(reorders, labels=['Reordered', 'Not Reordered'],
autopct='%1.1f%%')
plt.title('Percentage of Reordered Products')
plt.show()
```

Percentage of Reordered Products



```
orders_per_user = orders['user_id'].value_counts()

plt.figure(figsize=(10,6))
sns.histplot(orders_per_user, bins=30, kde=False)
plt.title('Number of Orders per User')
plt.show()
```



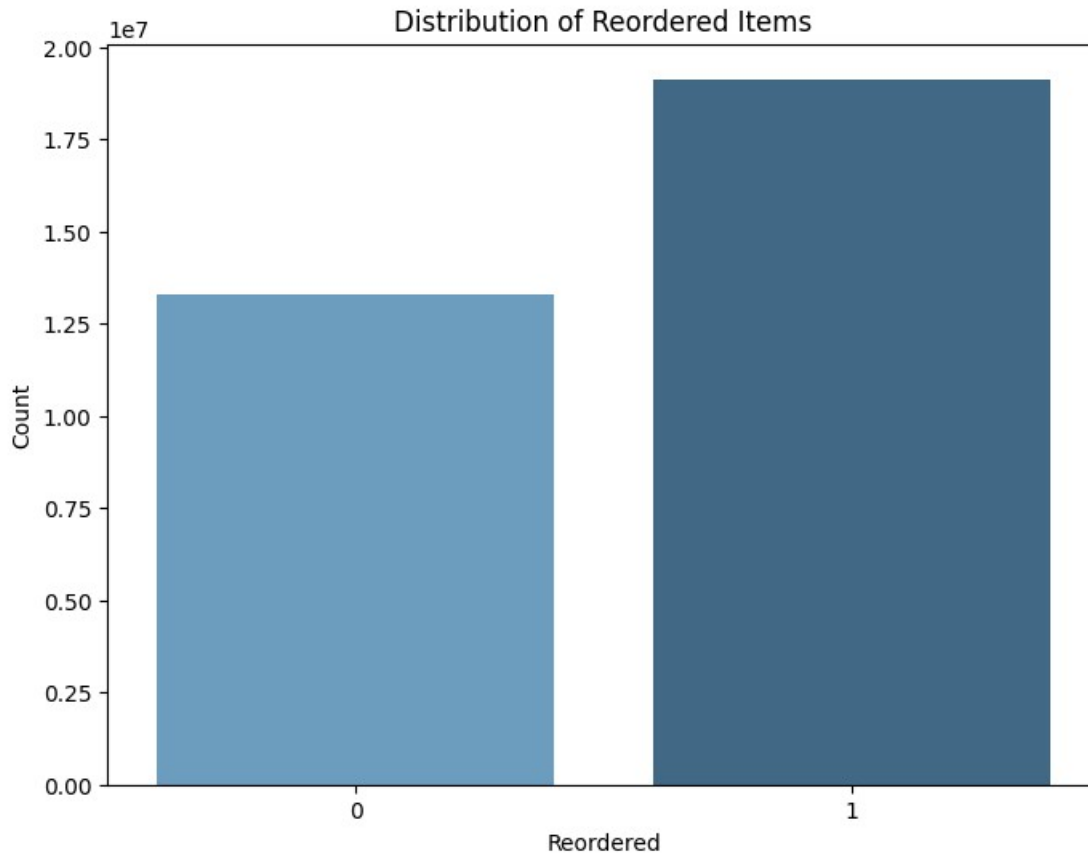
The tail on the right represents the regular customers, who orders frequently, while the peak on left represents the customers who orders less frequently.

*# Calculate the distribution of reordered values*

```
reorder_counts = order_products__prior['reordered'].value_counts()
```

*# Plot*

```
plt.figure(figsize=(8, 6))
sns.barplot(x=reorder_counts.index, y=reorder_counts.values,
palette="Blues_d")
plt.title('Distribution of Reordered Items')
plt.xlabel('Reordered')
plt.ylabel('Count')
plt.show()
```



## Section B: Data Pre-Processing

### Merging the data

```
prior_orders = pd.merge(order_products__prior, orders, on='order_id',  
how='left')
```

```
prior_orders.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32434489 entries, 0 to 32434488  
Data columns (total 10 columns):  
#   Column                                Dtype  
---  -  
0   order_id                             int64  
1   product_id                           int64  
2   add_to_cart_order                     int64  
3   reordered                             int64  
4   user_id                               int64  
5   eval_set                              object  
6   order_number                          int64  
7   order_dow                             int64  
8   order_hour_of_day                     int64  
9   days_since_prior_order                float64
```



```
dtypes: float64(1), int64(8), object(1)
memory usage: 2.4+ GB
```

### Downcasting the types of data to save memory

```
prior_orders['order_id'] = pd.to_numeric(prior_orders['order_id'],
downcast='integer')
prior_orders['product_id'] = pd.to_numeric(prior_orders['product_id'],
downcast='integer')
prior_orders['add_to_cart_order'] =
pd.to_numeric(prior_orders['add_to_cart_order'], downcast='integer')
prior_orders['reordered'] = pd.to_numeric(prior_orders['reordered'],
downcast='integer')
prior_orders['user_id'] = pd.to_numeric(prior_orders['user_id'],
downcast='integer')
prior_orders['order_number'] =
pd.to_numeric(prior_orders['order_number'], downcast='integer')
prior_orders['order_dow'] = pd.to_numeric(prior_orders['order_dow'],
downcast='integer')
prior_orders['order_hour_of_day'] =
pd.to_numeric(prior_orders['order_hour_of_day'], downcast='integer')
prior_orders['days_since_prior_order'] =
pd.to_numeric(prior_orders['days_since_prior_order'],
downcast='float')
```

### After downcasting:

```
prior_orders.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32434489 entries, 0 to 32434488
Data columns (total 10 columns):
```

#	Column	Dtype
0	order_id	int32
1	product_id	int32
2	add_to_cart_order	int16
3	reordered	int8
4	user_id	int32
5	eval_set	object
6	order_number	int8
7	order_dow	int8
8	order_hour_of_day	int8
9	days_since_prior_order	float32

```
dtypes: float32(1), int16(1), int32(3), int8(4), object(1)
memory usage: 928.0+ MB
```

### Checking and dealing with missing values

```
prior_orders.isnull().sum()
```

order_id	0
product_id	0
add_to_cart_order	0

```

reordered                0
user_id                  0
eval_set                 0
order_number             0
order_dow                0
order_hour_of_day        0
days_since_prior_order  2078068
dtype: int64

```

Imputation: only days\_since\_prior\_order has missing values, we will fill them with -1, which means that the order was made for the first time, 0 means that the order was made in the same after the previous one.

```

# Filling missing values in days_since_prior_order with -1
prior_orders['days_since_prior_order'].fillna(-1, inplace=True)

print(prior_orders.isnull().sum()/len(prior_orders))

```

```

order_id                0.0
product_id              0.0
add_to_cart_order       0.0
reordered               0.0
user_id                 0.0
eval_set                0.0
order_number            0.0
order_dow               0.0
order_hour_of_day       0.0
days_since_prior_order  0.0
dtype: float64

```

## Feature Engineering

We want to create features that describes:

1. each user's behaviour.
2. each product.
3. each user's behaviour towards the products.

```

# creating dataframe for the user's features.
userData = prior_orders[['user_id']].drop_duplicates()

# creating dataframe for the product's features.
prodData = prior_orders[['product_id']].drop_duplicates()

```

```

# creating dataframe for the user's behaviour towards the products.
prodUserData = prior_orders[['user_id',
'product_id']].drop_duplicates()

```

### Creating requested features:

Weekday Purchase Probability

```

# Creating a new feature: probability of a product being ordered on a
specific day of the week
total_orders = prior_orders.groupby('product_id')['order_id'].count()
day_of_week_orders = prior_orders.groupby(['product_id', 'order_dow'])
['order_id'].count()
weekday_purchase_prob = day_of_week_orders / total_orders
prodData['probability'] =
weekday_purchase_prob.reset_index(name='weekday_purchase_prob')
['weekday_purchase_prob']
merged_prior = prior_orders.copy()

```

### Product Purchase Frequency

```

# Creating a new feature: previous purchases of a product by a user
prodUserData['previous_purchases'] = merged_prior.groupby(['user_id',
'product_id'])['order_id'].transform('count')

```

### Days Since Last Order

```

# Creating a new feature: days since last order
merged_prior = merged_prior.sort_values(by=['user_id', 'product_id',
'order_number'])
merged_prior['order_gap'] = merged_prior.groupby(['user_id',
'product_id'])['order_number'].diff()
merged_prior['days_since_last_order'] =
merged_prior.groupby(['user_id', 'product_id'])
['days_since_prior_order'].cumsum()
prodUserData['days_since_last_order'] =
merged_prior['days_since_last_order'] -
merged_prior['days_since_prior_order']
merged_prior = merged_prior.drop(['order_gap'], axis=1)

```

### Creating another 7 features:

#### # 1. User's total orders

```

userData['total_orders'] = merged_prior.groupby('user_id')
['order_number'].transform('max')

```

#### # 2. Average basket size

```

userData['avg_basket_size'] = merged_prior.groupby('user_id')
['product_id'].transform('count') / userData['total_orders']

```

#### # 3. Average reorder ratio

```

userData['avg_reorder_ratio'] = merged_prior.groupby('user_id')
['reordered'].transform('mean')

```

#### # 4. Product purchase frequency

```

prodUserData['product_purchase_frequency'] =
merged_prior.groupby('product_id')['user_id'].transform('count') /
userData['user_id'].nunique()

```

```
# 5. Order_hour_of_day_period (Assuming 0-6 is night, 7-12 is morning,
13-18 is afternoon and 19-23 is evening)
```

```
bins = [0, 6, 12, 18, 23]
```

```
#Night=1, Morning=2, Afternoon=3, Evening=4
```

```
labels = ['0', '1', '2', '3']
```

```
userData['order_hour_of_day_period'] =
```

```
pd.cut(merged_prior['order_hour_of_day'], bins=bins, labels=labels,
include_lowest=True)
```

```
#prod ratio
```

```
prodData['prod_ratio'] = merged_prior.groupby('product_id')
```

```
['order_id'].transform('count') / merged_prior['order_id']
```

```
#average days between orders
```

```
userData['average_days_between_orders'] =
```

```
merged_prior.groupby('user_id')
```

```
['days_since_prior_order'].transform('mean')
```

```
userData
```

	user_id	total_orders	avg_basket_size	avg_reorder_ratio	
0	202279	8	11.250000	0.477778	\
9	205970	25	12.920000	0.588235	
17	178520	56	16.482143	0.884074	
30	156122	52	18.596154	0.840745	
56	22352	8	11.375000	0.087912	
...	...	...	...	...	
30480369	106087	3	8.333333	0.120000	
30557497	9295	3	3.333333	0.400000	
30756656	106586	3	5.000000	0.066667	
30944334	181902	3	20.333333	0.262295	
31084004	179441	3	7.000000	0.285714	

	order_hour_of_day_period	average_days_between_orders
0	1	19.322222
9	2	9.337461
17	1	6.631636
30	2	6.834540
56	1	22.857143
...	...	...
30480369	2	1.800000
30557497	3	20.400000
30756656	1	15.533334
30944334	3	3.901639
31084004	1	14.190476

```
[206209 rows x 6 columns]
```

```
prodData
```

	product_id	probability	prod_ratio
0	33120	0.108531	9.700000e+03
1	28985	0.215983	3.372450e+04
2	9327	0.146328	3.149500e+03
3	45918	0.141469	3.720000e+02
4	30035	0.159287	2.845000e+02
...	...	...	...
30670613	7726	NaN	6.182264e-07
30954037	43492	NaN	6.125389e-07
30974732	33097	NaN	3.060724e-07
32168896	38977	NaN	2.947113e-07
32242079	23624	NaN	2.940482e-07

[49677 rows x 3 columns]

prodUserData

	user_id	product_id	previous_purchases
days_since_last_order			
0	202279	33120	5
1.0 \			
1	202279	28985	5
19.0			
2	202279	9327	1
0.0			
3	202279	45918	5
20.0			
4	202279	30035	3
0.0			
...	...	...	...
..			
32434480	25247	45309	1
0.0			
32434481	25247	21162	1
0.0			
32434483	25247	35211	1
0.0			
32434485	25247	11352	1
0.0			
32434486	25247	4600	1
0.0			

	product_purchase_frequency
0	0.094079
1	0.327090
2	0.030547
3	0.003608
4	0.002759
...	...
32434480	0.005941

```

32434481          0.026977
32434483          0.000548
32434485          0.034281
32434486          0.012730

```

[13307953 rows x 5 columns]

**Transformation - Normalizing, and Splitting the new Features to numerical and catagorical.**

```

def normalize(df,k):
    for col in df.columns[k:]:
        if len(df[col].unique()) <= 25 or col.endswith('dow') or
col.endswith('hour') or col.endswith('period') or
col.endswith('preference'):
            df[col] = df[col].astype('category')
        else:
            df[col] = df[col].astype('float32')

```

*#call normalize with userData, prodData, prodUserData*

```

normalize(userData, 1)
normalize(prodData, 1)
normalize(prodUserData, 2)

```

Normalized

## Section C: Future Order Prediction

**Creating Train and Test Dataframes**

*#merge orders with target*

```

final = pd.merge(target, orders, on='order_id', how='left')

```

final

	order_id	product_id	was_in_order	user_id	eval_set
order_number					
0	1187899	196.0	1	1	test
11 \					
1	2757217	196.0	1	67	test
25					
2	632715	196.0	1	676	train
12					
3	1167274	196.0	1	760	test
5					
4	3347074	196.0	1	804	train
16					
...	...	...	...	...	...
...					
2959120	3421070	4987.0	0	139822	test
15					
2959121	3421070	8230.0	0	139822	test
15					

2959122	3421070	39468.0	0	139822	test
15					
2959123	3421070	39139.0	0	139822	test
15					
2959124	3421070	48208.0	0	139822	test
15					

	order_dow	order_hour_of_day	days_since_prior_order
0	4	8	14.0
1	0	11	5.0
2	0	13	26.0
3	4	10	8.0
4	3	21	5.0
...	...	...	...
2959120	6	10	8.0
2959121	6	10	8.0
2959122	6	10	8.0
2959123	6	10	8.0
2959124	6	10	8.0

[2959125 rows x 9 columns]

```
final = final[['user_id', 'eval_set',
'order_id', 'product_id', 'was_in_order']]
```

*#merging prodData with order\_train\_test*

```
order_train_test = pd.merge(final,userData, on='user_id', how='left')
```

*#merging prodData with order\_train\_test*

```
order_train_test = pd.merge(order_train_test,prodData,
on='product_id', how='left')
```

*#merging prodUserData with order\_train\_test*

```
order_train_test = pd.merge(order_train_test,prodUserData,
on=['user_id', 'product_id'], how='left')
```

order\_train\_test

	user_id	eval_set	order_id	product_id	was_in_order
total_orders					
0	1	test	1187899	196.0	1
10.0 \					
1	67	test	2757217	196.0	1
24.0					
2	676	train	632715	196.0	1
11.0					
3	760	test	1167274	196.0	1
4.0					
4	804	train	3347074	196.0	1
15.0					
...	...	...	...	...	...

...					
2959120	139822	test	3421070	4987.0	0
14.0					
2959121	139822	test	3421070	8230.0	0
14.0					
2959122	139822	test	3421070	39468.0	0
14.0					
2959123	139822	test	3421070	39139.0	0
14.0					
2959124	139822	test	3421070	48208.0	0
14.0					

	avg_basket_size	avg_reorder_ratio	order_hour_of_day_period
0	5.900000	0.694915	
2 \			
1	3.375000	0.716049	2
2	1.727273	0.631579	1
3	2.000000	0.375000	1
4	7.466667	0.732143	1
...	...	...	...
2959120	2.142857	0.633333	2
2959121	2.142857	0.633333	2
2959122	2.142857	0.633333	2
2959123	2.142857	0.633333	2
2959124	2.142857	0.633333	2

	average_days_between_orders	probability	prod_ratio	
0	18.457626	0.500000	2753.153809	\
1	7.283951	0.500000	2753.153809	
2	19.315790	0.500000	2753.153809	
3	7.375000	0.500000	2753.153809	
4	16.178572	0.500000	2753.153809	
...	...	...	...	
2959120	11.066667	0.033333	0.084243	
2959121	11.066667	0.153846	42.982250	
2959122	11.066667	0.250000	5.676349	
2959123	11.066667	0.154472	0.015088	
2959124	11.066667	NaN	0.000065	



	previous_purchases	days_since_last_order
product_purchase_frequency		
0	10.0	64.0
0.173567		
1	19.0	95.0
0.173567		
2	7.0	14.0
0.173567		
3	1.0	0.0
0.173567		
4	NaN	NaN
NaN		
...	...	...
...		
2959120	NaN	NaN
NaN		
2959121	NaN	NaN
NaN		
2959122	NaN	NaN
NaN		
2959123	NaN	NaN
NaN		
2959124	NaN	NaN
NaN		

[2959125 rows x 15 columns]

*#display null ratio*

```
print(order_train_test.isnull().sum()/len(order_train_test))
```

user_id	0.000000
eval_set	0.000000
order_id	0.000000
product_id	0.000000
was_in_order	0.000000
total_orders	0.000000
avg_basket_size	0.000000
avg_reorder_ratio	0.000000
order_hour_of_day_period	0.000000
average_days_between_orders	0.000000
probability	0.069912
prod_ratio	0.000019
previous_purchases	0.703579
days_since_last_order	0.703579
product_purchase_frequency	0.703579
dtype: float64	

note : it's better to drop the columns that has >50% nulls, but because 2 were given in the homework, we will not delete any.

### Distribution of numerical values - before filling NaNs

```
# order_train_test.iloc[:, -10:].hist(bins=50, figsize=(15, 15))

#interpolate using group by user_id lambda x method ffill
order_train_test = order_train_test.groupby('user_id').apply(lambda x:
x.interpolate(method='ffill'))

#filling last nulls with 0
for col in order_train_test.columns[-5:]:
    order_train_test[col].fillna(0, inplace=True)

print(order_train_test.isnull().sum()/len(order_train_test))

user_id                0.0
eval_set               0.0
order_id              0.0
product_id            0.0
was_in_order          0.0
total_orders          0.0
avg_basket_size       0.0
avg_reorder_ratio     0.0
order_hour_of_day_period 0.0
average_days_between_orders 0.0
probability           0.0
prod_ratio            0.0
previous_purchases    0.0
days_since_last_order 0.0
product_purchase_frequency 0.0
dtype: float64
```

### Distribution of numerical values - after filling NaNs

```
# order_train_test.iloc[:, -10:].hist(bins=50, figsize=(15, 15))
```

Probability: The statistics remain roughly the same except that the mean becomes negative (approaches zero). The negative skewness suggests a longer tail on the left side of the distribution.

Previous Purchases: Similarly, the mean is near zero and standard deviation is near 1, which means data are standardized. The positive skewness indicates a longer tail on the right side, suggesting more customers with lower previous purchases.

User Total Orders: The data distribution is similar to that of Previous Purchases, but it is negatively skewed now, indicating a longer tail on the left side.

Avg Days Between Orders: The skewness is negative, suggesting that there are more instances of shorter average days between orders. The low kurtosis suggests the distribution is platykurtic (less outliers or extreme values).

Avg Basket Size: The skewness is slightly negative suggesting a longer tail on the left side of the distribution. The positive kurtosis indicates a leptokurtic distribution, suggesting there may be more extreme values or outliers.

Avg Reorder Ratio: The distribution now shows a negative skewness, indicating a longer tail on the left side. The positive kurtosis suggests more outliers in the data.

Product Purchase Frequency: The skewness is positive and quite large, indicating a longer tail on the right side. The kurtosis is also quite large, suggesting that the distribution has heavy tails and more outliers.

Product Reorder Ratio: The distribution is negatively skewed, suggesting a longer tail on the left side. The kurtosis is greater than zero, indicating a leptokurtic distribution with more outliers or extreme values.

User Product Reorder Ratio: The negative skewness suggests a longer tail on the left side of the distribution. The kurtosis being negative suggests a platykurtic distribution, with fewer outliers or extreme values than a normal distribution.

we can reject the null hypothesis of the data being normally distributed.

This suggests that although the standardization has adjusted the scale of the data, it hasn't made the data follow a normal distribution.

## One-Hot Encoding

```
import pandas as pd

def analyze_and_encode(df):
    for col in df.columns:
        if df[col].dtype == 'category':
            print("column: ", col, "unique value counts: ",
df[col].value_counts())

    # Applying one-hot encoding to 'order_hour_of_day_period'
    df_encoded = pd.get_dummies(df, columns=[col])
    #drop the original column
    return df_encoded

#call analyze_and_encode function
order_train_test = analyze_and_encode(order_train_test)

column:  order_hour_of_day_period unique value counts:
order_hour_of_day_period
2      1366584
1      1137108
3       376152
0        79281
Name: count, dtype: int64

order_train_test.info()

<class 'pandas.core.frame.DataFrame'>
MultiIndex: 2959125 entries, (1, 0) to (206209, 1510784)
Data columns (total 18 columns):
#   Column                                Dtype
```

```

---  -----
0   user_id          int64
1   eval_set         object
2   order_id         int64
3   product_id       float64
4   was_in_order     int64
5   total_orders     float32
6   avg_basket_size  float32
7   avg_reorder_ratio float32
8   average_days_between_orders float32
9   probability      float32
10  prod_ratio       float32
11  previous_purchases float32
12  days_since_last_order float32
13  product_purchase_frequency float32
14  order_hour_of_day_period_0 bool
15  order_hour_of_day_period_1 bool
16  order_hour_of_day_period_2 bool
17  order_hour_of_day_period_3 bool
dtypes: bool(4), float32(9), float64(1), int64(3), object(1)
memory usage: 340.4+ MB

```

order\_train\_test

	user_id	eval_set	order_id	product_id	was_in_order	
user_id						
1	0	1	test	1187899	196.0	1
\	2272	1	test	1187899	26088.0	1
	2408	1	test	1187899	26405.0	1
	2513	1	test	1187899	10258.0	1
	18097	1	test	1187899	13032.0	1
...	...	...	...	...	...	
206209	1510780	206209	train	272231	24838.0	0
	1510781	206209	train	272231	14678.0	0
	1510782	206209	train	272231	31717.0	0
	1510783	206209	train	272231	49286.0	0
	1510784	206209	train	272231	1184.0	0

user_id		total_orders	avg_basket_size	avg_reorder_ratio
1	0	10.0	5.900000	0.694915 \
	2272	10.0	5.900000	0.694915
	2408	10.0	5.900000	0.694915
	2513	10.0	5.900000	0.694915
	18097	10.0	5.900000	0.694915
...		...	...	...
206209	1510780	13.0	9.923077	0.472868
	1510781	13.0	9.923077	0.472868
	1510782	13.0	9.923077	0.472868
	1510783	13.0	9.923077	0.472868
	1510784	13.0	9.923077	0.472868

prod_ratio		average_days_between_orders	probability
user_id			
1	0	18.457626	0.500000
2753.153809 \	2272	18.457626	0.146739
6.403553	2408	18.457626	0.111111
0.267519	2513	18.457626	0.196078
4.334076	18097	18.457626	0.500000
9.692507		...	...
...		...	...
206209	1510780	18.131783	0.166667
16523.000000	1510781	18.131783	0.046512
145.525177	1510782	18.131783	0.154472
6952.399902	1510783	18.131783	0.154472
0.000055	1510784	18.131783	0.140401
0.001093			

user_id		previous_purchases	days_since_last_order
1	0	10.0	64.0 \
	2272	2.0	-1.0
	2408	2.0	-1.0
	2513	9.0	65.0
	18097	3.0	15.0

...		...	...
206209	1510780	10.0	39.0
	1510781	10.0	39.0
	1510782	10.0	39.0
	1510783	10.0	39.0
	1510784	10.0	39.0

	product_purchase_frequency
order_hour_of_day_period_0	
user_id	

1	0	0.173567
False	\	
	2272	0.012235
False		
	2408	0.005887
False		
	2513	0.009437
False		
	18097	0.018190
False		
...	...	..
.		
206209	1510780	0.014917
False		
	1510781	0.014917
False		
	1510782	0.014917
False		
	1510783	0.014917
False		
	1510784	0.014917
False		

	order_hour_of_day_period_1
order_hour_of_day_period_2	
user_id	

1	0	False
True	\	
	2272	False
True		
	2408	False
True		
	2513	False
True		
	18097	False
True		
...	...	..
.		

206209	1510780	False
True		
	1510781	False
True		
	1510782	False
True		
	1510783	False
True		
	1510784	False
True		

		order_hour_of_day_period_3
user_id		
1	0	False
	2272	False
	2408	False
	2513	False
	18097	False
...		...
206209	1510780	False
	1510781	False
	1510782	False
	1510783	False
	1510784	False

[2959125 rows x 18 columns]

### Splitting that data to train and test sets and splitting the feature from the labels

```
# #split train into X and y
# X = order_train_test.drop(['was_in_order'], axis=1)
# y = order_train_test['was_in_order']

# #import smote
# from imblearn.over_sampling import SMOTE
# sm = SMOTE(random_state=42)
# X_res, y_res = sm.fit_resample(X, y)

# #do the same

train_data = order_train_test[order_train_test['eval_set'] == 'train']
train_data = train_data.set_index(['user_id', 'product_id'])
train_data.drop(['eval_set', 'order_id'], axis=1, inplace=True)
train_data = train_data.reset_index()

test_data = order_train_test[order_train_test['eval_set'] == 'test']
test_data = test_data.set_index(['user_id', 'product_id'])
test_data.drop(['eval_set', 'order_id'], axis=1, inplace=True)
test_data = test_data.reset_index()
```

```

#split train into X and y
X = train_data.drop(['was_in_order'], axis=1)
y = train_data['was_in_order']

#import smote
from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X, y)

X_test = test_data.drop(['was_in_order'], axis=1)

y_res.value_counts()

was_in_order
1    1102152
0    1102152
Name: count, dtype: int64

Synthecizing the data so the label is split 50/50

y_train = train_data['was_in_order']
X_train = train_data.drop(['was_in_order'], axis=1)

y_test = test_data['was_in_order']
X_test = test_data.drop(['was_in_order'], axis=1)

#split train into X and y

#import smote
from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)
X_train, y_train = sm.fit_resample(X_train, y_train)
X_test, y_test = sm.fit_resample(X_test, y_test)

X_test.drop(['user_id', 'product_id'], axis=1, inplace=True)
X_train.drop(['user_id', 'product_id'], axis=1, inplace=True)

#value counts
y_train.value_counts()

was_in_order
1    1102152
0    1102152
Name: count, dtype: int64

#value counts
y_test.value_counts()

was_in_order
1    472356

```



```
0    472356
Name: count, dtype: int64
```

**we will use 3 different metrics to evaluate the models: (ref: ML Example, Moodle)**

1. Accuracy: refers to the ability to correctly predict both positive and negative observations.
2. Sensitivity: refers to the ability to correctly predict positive observations.
3. Specificity: refers to the ability to correctly predict negative observations.

```
def get_metrics(y_test, y_pred):

    # actual_pos = y_test == 1
    # actual_neg = y_test == 0

    # # get confusion matrix
    # mat = metrics.confusion_matrix(y_test, y_pred)
    # true_neg, false_pos, false_neg, true_pos = mat.ravel()

    mat = metrics.confusion_matrix(y_test, y_pred)

    # handle case when confusion matrix is a single value
    if mat.shape == (1, 1):
        total_samples = mat.item()
        actual_pos = false_pos = false_neg = true_neg = total_samples
    else:
        actual_pos, false_pos, false_neg, true_neg = mat.ravel()

    # calculate sensitivity and specificity
    sensitivity = round(actual_pos / (actual_pos+false_neg), 3)
    specificity = round(true_neg / (true_neg + false_pos), 3)
    accuracy = accuracy_score(y_test, y_pred)
    F1_score = f1_score(y_test, y_pred)
    Precision = precision_score(y_test, y_pred)
    Recall = recall_score(y_test, y_pred)

    return sensitivity, specificity, accuracy, F1_score, Precision,
Recall
```

**We are using 3 different models to predict the future orders:**

## Ada Boost Classifier

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import AdaBoostClassifier

# Set the parameters by cross-validation
param_grid = {'n_estimators': [10, 25], 'learning_rate': [0.01, 0.3]}
```

```

# Instantiate the grid search model
grid_search_ab = GridSearchCV(estimator = AdaBoostClassifier(),
param_grid = param_grid, cv = 3)

# Fit the grid search to the data
grid_search_ab.fit(X_train, y_train)

# Print the best parameters
print("Best parameters: ", grid_search_ab.best_params_)
y_pred = grid_search_ab.predict(X_test)

# Display metrics
sensitivity, specificity, accuracy, F1_score, Precision, Recall =
get_metrics(y_test, y_pred)
ab_df = pd.DataFrame([accuracy, sensitivity, specificity, F1_score,
Precision, Recall]).T
ab_df = ab_df.rename(index={0: 'Ada Boost Classifier'}, columns={0:
'Accuracy', 1: 'Sensitivity', 2: 'Specificity', 3: 'F1_score', 4:
'Precision', 5: 'Recall'})
ab_df

Best parameters: {'learning_rate': 0.3, 'n_estimators': 25}

                Accuracy  Sensitivity  Specificity  F1_score
Precision
Ada Boost Classifier  0.680449          0.69          0.672  0.688372
0.671717  \


                Recall
Ada Boost Classifier  0.705874

comparison_df = pd.DataFrame(columns=['Model', 'Accuracy',
'Sensitivity', 'Specificity', 'F1_score', 'Precision', 'Recall'])
row = pd.Series({'Model': 'Ada Boost Classifier', 'Accuracy':
accuracy, 'Sensitivity': sensitivity, 'Specificity': specificity,
'F1_score': F1_score, 'Precision': Precision, 'Recall': Recall})
comparison_df.loc[0] = row
del accuracy, sensitivity, specificity, F1_score, Precision, Recall

cm = confusion_matrix(y_test, y_pred)
# Create the ConfusionMatrixDisplay object
cmd = ConfusionMatrixDisplay(cm)

# Plot the confusion matrix
cmd.plot()
plt.show()

```



## Decision Tree

*# defining the model*

```
clf = DecisionTreeClassifier()
```

*# define grid search*

```
paramGrid = {
```

```
    #fill more on stronger PCs
```

```
    "max_depth": [3, 5],
```

```
    "min_samples_split": [2],
```

```
    "min_samples_leaf": [1],
```

```
    "criterion": ["gini", "entropy"],
```

```
}
```

```
grid_search_dt = GridSearchCV(estimator=clf, param_grid=paramGrid,  
cv=3, scoring='accuracy')
```

*# fit the model*

```
grid_search_dt.fit(X_train, y_train)
```

*# After the fitting process, print the best parameters and the best score*

```
print("Best Parameters: ", grid_search_dt.best_params_)
```

*# Use the grid search (with best parameters) to make predictions on*

*the test set*

```
best = grid_search_dt.best_estimator_  
y_pred = best.predict(X_test)  
print(y_pred)
```

```
Best Parameters: {'criterion': 'gini', 'max_depth': 5,  
'min_samples_leaf': 1, 'min_samples_split': 2}  
[1 0 0 ... 1 0 1]
```

*# display metrics*

```
sensitivity, specificity, accuracy, F1_score, Precision, Recall =  
get_metrics(y_test, y_pred)  
dt_df = pd.DataFrame([accuracy, sensitivity, specificity, F1_score,  
Precision, Recall]).T  
dt_df = dt_df.rename(index={0: 'Decision Tree Classifier'},  
columns={0: 'Accuracy', 1: 'Sensitivity', 2: 'Specificity' , 3:  
'F1_score', 4: 'Precision', 5: 'Recall'})  
dt_df
```

	Accuracy	Sensitivity	Specificity	F1_score
Decision Tree Classifier	0.715672	0.739	0.696	0.72894


	Precision	Recall
Decision Tree Classifier	0.696441	0.76462

```
row = pd.Series({'Model': 'Decision Tree Classifier', 'Accuracy':  
accuracy, 'Sensitivity': sensitivity, 'Specificity': specificity,  
'F1_score': F1_score, 'Precision': Precision, 'Recall': Recall})  
comparison_df.loc[1] = row  
del accuracy, sensitivity, specificity, F1_score, Precision, Recall
```

```
cm = confusion_matrix(y_test, y_pred)  
# Create the ConfusionMatrixDisplay object  
cmd = ConfusionMatrixDisplay(cm)
```

*# Plot the confusion matrix*

```
cmd.plot()  
plt.show()
```



## XGB

```
paramGrid = {"max_depth":[5,8], "colsample_bytree":[0.3,0.5]}
```

```
# Define the XGBoost classifier
```

```
xgbc = XGBClassifier(objective='binary:logistic',  
eval_metric='logloss', n_estimators=10)
```

```
# Apply GridSearchCV for parameter tuning
```

```
gridsearch = GridSearchCV(xgbc, paramGrid, cv=3, verbose=2, n_jobs=1)
```

```
model = gridsearch.fit(X_train, y_train)
```

```
print("The best parameters are: \n", gridsearch.best_params_)
```

```
model = gridsearch.best_estimator_
```

```
Fitting 3 folds for each of 4 candidates, totalling 12 fits
```

```
[CV] END .....colsample_bytree=0.3, max_depth=5; total  
time= 2.4s
```

```
[CV] END .....colsample_bytree=0.3, max_depth=5; total  
time= 2.4s
```

```
[CV] END .....colsample_bytree=0.3, max_depth=5; total  
time= 2.3s
```

```
[CV] END .....colsample_bytree=0.3, max_depth=8; total  
time= 3.0s
```

```
[CV] END .....colsample_bytree=0.3, max_depth=8; total
time= 3.0s
[CV] END .....colsample_bytree=0.3, max_depth=8; total
time= 2.9s
[CV] END .....colsample_bytree=0.5, max_depth=5; total
time= 2.7s
[CV] END .....colsample_bytree=0.5, max_depth=5; total
time= 2.7s
[CV] END .....colsample_bytree=0.5, max_depth=5; total
time= 2.6s
[CV] END .....colsample_bytree=0.5, max_depth=8; total
time= 3.7s
[CV] END .....colsample_bytree=0.5, max_depth=8; total
time= 3.6s
[CV] END .....colsample_bytree=0.5, max_depth=8; total
time= 3.5s
The best parameters are: /n {'colsample_bytree': 0.5, 'max_depth': 8}
```

```
y_pred = model.predict(X_test)
print(y_pred)
```

```
# display metrics
```

```
sensitivity, specificity, accuracy, F1_score, Precision, Recall =
get_metrics(y_test, y_pred)
gb_df = pd.DataFrame([accuracy, sensitivity, specificity, F1_score,
Precision, Recall]).T
gb_df = gb_df.rename(index={0: 'Gradient Boosting Classifier'},
columns={0: 'Accuracy', 1: 'Sensitivity', 2: 'Specificity' , 3:
'F1_score', 4: 'Precision', 5: 'Recall'})
gb_df
```

```
[1 1 1 ... 1 1 1]
```

	Accuracy	Sensitivity	Specificity
F1_score			
Gradient Boosting Classifier	0.768668	0.802	0.742
0.780663 \			


	Precision	Recall
Gradient Boosting Classifier	0.742181	0.823354

```
row = pd.Series({'Model': 'Gardient Boosting Classifier', 'Accuracy':
accuracy, 'Sensitivity': sensitivity, 'Specificity': specificity,
'F1_score': F1_score, 'Precision': Precision, 'Recall': Recall})
comparison_df.loc[2] = row
```

```
cm = confusion_matrix(y_test, y_pred)
# Create the ConfusionMatrixDisplay object
cmd = ConfusionMatrixDisplay(cm)
```


```
# Plot the confusion matrix
```

```
cmd.plot()  
plt.show()
```



#### Comparing the results of the models

```
comparison_df.plot(kind='bar', x='Model', y=['Accuracy',  
'Sensitivity', 'Specificity', 'F1_score', 'Precision', 'Recall'],  
figsize=(10,5))  
plt.title('Comparison of Models')  
plt.ylabel('Score')  
plt.show()
```



## Section D: Clustering

Creating new features that describes customers buying behavior

```
clus_prior_orders = pd.merge(order_products__prior, orders,
on='order_id', how='left')
clus_prior_orders = pd.merge(clus_prior_orders, products,
on='product_id', how='left')
clus_prior_orders = clus_prior_orders[['user_id', 'product_id',
'aisle_id', 'department_id', 'order_id', 'order_dow',
'order_hour_of_day', 'days_since_prior_order']]

#calculate favorite day of week for each user
user_dow = clus_prior_orders.groupby(['user_id',
'order_dow']).size().reset_index(name='count')
user_dow = user_dow.sort_values(by=['user_id', 'count'],
ascending=False)
user_dow = user_dow.drop_duplicates(subset=['user_id'], keep='first')
user_dow = user_dow[['user_id', 'order_dow']]
user_dow = user_dow.rename(columns={'order_dow': 'favorite_dow'})
clus_prior_orders = pd.merge(clus_prior_orders, user_dow,
on='user_id', how='left')
```



```

#favourite department for each user
user_department = clus_prior_orders.groupby(['user_id',
'department_id']).size().reset_index(name='count')
user_department = user_department.sort_values(by=['user_id', 'count'],
ascending=False)
user_department = user_department.drop_duplicates(subset=['user_id'],
keep='first')
user_department = user_department[['user_id', 'department_id']]
user_department = user_department.rename(columns={'department_id':
'favorite_department'})
clus_prior_orders = pd.merge(clus_prior_orders, user_department,
on='user_id', how='left')

```

```

clus_prior_orders = clus_prior_orders[['user_id',
'favorite_dow', 'favorite_department']]
clus_prior_orders =
clus_prior_orders.drop_duplicates(subset=['user_id'], keep='first')

clus_prior_orders = pd.merge(clus_prior_orders, userData,
on='user_id', how='left')

```

clus\_prior\_orders

	user_id	favorite_dow	favorite_department	total_orders	
0	202279	5	13	8.0	\
1	205970	6	4	25.0	
2	178520	1	16	56.0	
3	156122	3	16	52.0	
4	22352	1	4	8.0	
...	...	...	...	...	
206204	106087	6	1	3.0	
206205	9295	6	16	3.0	
206206	106586	5	13	3.0	
206207	181902	0	4	3.0	
206208	179441	0	4	3.0	

	avg_basket_size	avg_reorder_ratio	order_hour_of_day_period	
0	11.250000	0.477778	1	\
1	12.920000	0.588235	2	
2	16.482143	0.884074	1	
3	18.596153	0.840745	2	
4	11.375000	0.087912	1	
...	...	...	...	
206204	8.333333	0.120000	2	
206205	3.333333	0.400000	3	
206206	5.000000	0.066667	1	
206207	20.333334	0.262295	3	
206208	7.000000	0.285714	1	

	average_days_between_orders
0	19.322222
1	9.337461
2	6.631636
3	6.834540
4	22.857143
...	...
206204	1.800000
206205	20.400000
206206	15.533334
206207	3.901639
206208	14.190476

[206209 rows x 8 columns]

```
clus_prior_orders = clus_prior_orders.drop(columns=['user_id'])
```

Might add more features

```
#import Kmeans and cdists
```

```
from sklearn.cluster import KMeans
```

```
from scipy.spatial.distance import cdist
```

```
def elbow_met(df):
```

```
    # k means determine k
```

```
    distortions = []
```

```
    K = range(1,10)
```

```
    for k in K:
```

```
        kmeanModel =
```

```
        KMeans(n_clusters=k,n_init=10).fit(clus_prior_orders)
```

```
        kmeanModel.fit(clus_prior_orders)
```

```
        distortions.append(sum(np.min(cdist(clus_prior_orders,
kmeanModel.cluster_centers_, 'euclidean'), axis=1)) /
clus_prior_orders.shape[0])
```

```
    # Plot the elbow
```

```
    plt.plot(K, distortions, 'bx-')
```

```
    plt.xlabel('k')
```

```
    plt.ylabel('Within groups sum of squares')
```

```
    plt.title('The Elbow Method showing the optimal k')
```

```
    plt.show()
```

```
clus_prior_orders
```

	favorite_dow	favorite_department	total_orders
avg_basket_size			
0	5	13	8.0
11.250000 \			
1	6	4	25.0
12.920000			

2	1	16	56.0
16.482143			
3	3	16	52.0
18.596153			
4	1	4	8.0
11.375000			
...	...	...	...
206204	6	1	3.0
8.333333			
206205	6	16	3.0
3.333333			
206206	5	13	3.0
5.000000			
206207	0	4	3.0
20.333334			
206208	0	4	3.0
7.000000			

	avg_reorder_ratio	order_hour_of_day_period	
0	0.477778	1	\
1	0.588235	2	
2	0.884074	1	
3	0.840745	2	
4	0.087912	1	
...	...	...	
206204	0.120000	2	
206205	0.400000	3	
206206	0.066667	1	
206207	0.262295	3	
206208	0.285714	1	

	average_days_between_orders
0	19.322222
1	9.337461
2	6.631636
3	6.834540
4	22.857143
...	...
206204	1.800000
206205	20.400000
206206	15.533334
206207	3.901639
206208	14.190476

[206209 rows x 7 columns]

elbow\_met(clus\_prior\_orders)

```
-----  
-----  
ValueError                                Traceback (most recent call  
last)
```

```
Cell In[70], line 1  
----> 1 elbow_met(clus_prior_orders)
```

```
Cell In[68], line 13, in elbow_met(df)  
     11     kmeanModel =  
KMeans(n_clusters=k,n_init=10).fit(clus_prior_orders)  
     12     kmeanModel.fit(clus_prior_orders)  
--> 13     distortions.append(sum(np.min(cdist(clus_prior_orders,  
kmeanModel.cluster_centers_, 'euclidean'), axis=1)) /  
clus_prior_orders.shape[0])  
     15 # Plot the elbow  
     16 plt.plot(K, distortions, 'bx-')
```

```
File  
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/  
site-packages/scipy/spatial/distance.py:2939, in cdist(XA, XB, metric,  
out, **kwargs)  
    2937 if metric_info is not None:  
    2938     cdist_fn = metric_info.cdist_func  
-> 2939     return cdist_fn(XA, XB, out=out, **kwargs)  
    2940 elif mstr.startswith("test"):  
    2941     metric_info = _TEST_METRICS.get(mstr, None)
```

```
ValueError: Unsupported dtype object
```

```
kmeans = KMeans(n_clusters=2,  
random_state=0,n_init=10).fit(clus_prior_orders)
```


```
centroids = kmeans.cluster_centers_  
print(centroids)
```

```
[[5.16770000e+04 2.52050739e+00 7.34059969e+00]  
 [1.54781500e+05 2.52199191e+00 7.37228747e+00]]
```

```
#run kmeans with 2 and plot
```

```
kmeans = KMeans(n_clusters=2,  
random_state=0,n_init=10).fit(clus_prior_orders)  
plt.scatter(clus_prior_orders.iloc[:,0], clus_prior_orders.iloc[:,1],  
c=kmeans.labels_, cmap='rainbow')
```

```
<matplotlib.collections.PathCollection at 0x13f05b79d50>
```



## BONUS

### PCA - 2 features

*#apply pca for User data*

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=2)
```

```
principalComponents = pca.fit_transform(cluster_prior)
```

```
principalDf = pd.DataFrame(data = principalComponents
```

```
, columns = ['principal component 1', 'principal  
component 2'])
```

```
principalDf
```

```
-----  
-----
```

```
NameError
```

```
Traceback (most recent call
```

```
last)
```

```
Cell In[88], line 4
```

```
2 from sklearn.decomposition import PCA
```

```
3 pca = PCA(n_components=2)
```

```
----> 4 principalComponents = pca.fit_transform(cluster_prior)
```

```
5 principalDf = pd.DataFrame(data = principalComponents
```

```
6
```

```
7          , columns = ['principal component 1',  
'principal component 2'])  
9 principalDf
```

NameError: name 'cluster\_prior' is not defined

*#plot PCA*

```
plt.scatter(principalDf.iloc[:, 0], principalDf.iloc[:, 1], c=labels,  
s=50, cmap='viridis')
```

<matplotlib.collections.PathCollection at 0x45600e3d0>

